



# Grid Dynamics

Scalable eCommerce Platform Solutions

## *Not a relational databases*

**NOSQL** Redis  
Cassandra  
Voldemort  
Key/Value EC2  
MapReduce  
HBase  
CAP  
BerkeleyDB  
Neo4J  
SimpleDB  
Cloud  
Riak  
memcached  
MongoDB  
Azure  
CouchDB  
Hadoop

by Roman Nikolaenko

# Topics

I part:


- Why RDBMS is not a "silver bullet"?
- Do we have any alternatives?
- What does CAP theorem mean?
- Any use cases for noSQL?
- What was our use case?
- HBase
- MongoDB

II part:

- MongoDB Demo

# Why RDBMS is not a "silver bullet"?

ACID\*  
Transactions  
Data consistency



\*ACID (*atomicity, consistency, isolation, durability*)

# Why RDBMS is not a "silver bullet"?

Amount of Data **is growing** rapidly.

Companies:

- Facebook
- Twitter
- MySpace
- Google
- Amazon

Store **petabytes** of data, make complex **analysis**

# Why RDBMS is not a "silver bullet"?

Amount of Data is **growing** rapidly.

Companies:

- Facebook
- Twitter
- MySpace
- Google
- Amazon


WHAT  
COULD  
WE DO?

Store **petabytes** of data with complex **analysis**

# Why RDBMS is not a "silver bullet"?

ACID\*  
Transactions  
Data consistency

+CPU  
+RAM  
+HDD  
...  
**PRICE**




\*ACID (*atomicity, consistency, isolation, durability*)

# Why RDBMS is not a "silver bullet"?

ACID ?  
Transactions ?  
Data consistency ?

Fault Tolerance ?  
Data replication ?  
Data sharding ?



# Why RDBMS is not a "silver bullet"?

## Scalability issues

We need additional stuff to maintain cluster



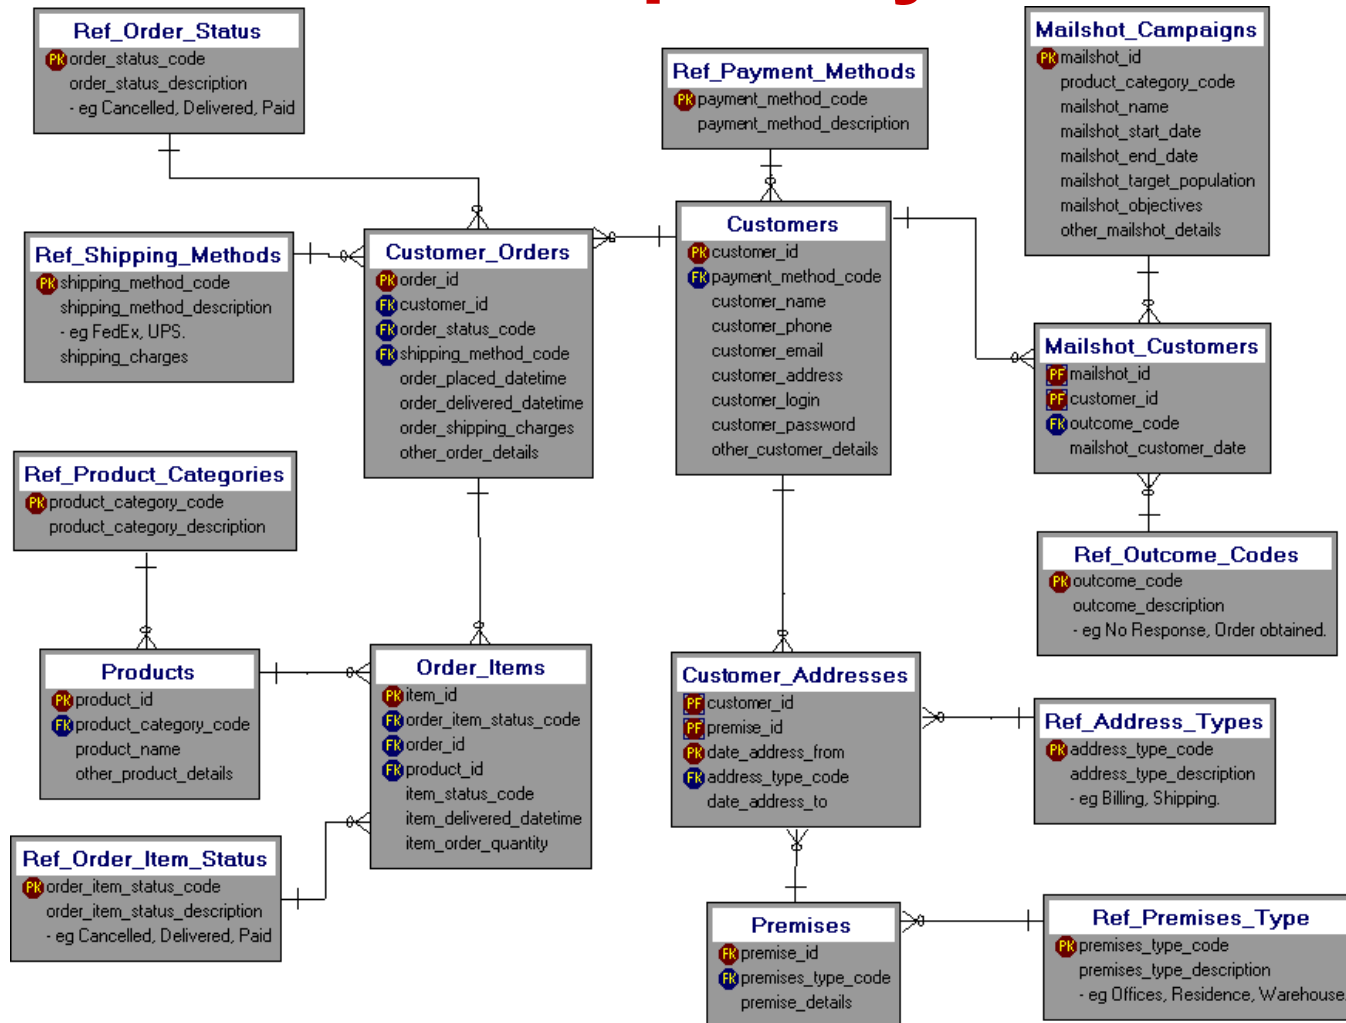
# Why RDBMS is not a "silver bullet"?

**Scalability issues**

**Complexity**

# Why RDBMS is not a "silver bullet"?

## Complexity



# Why RDBMS is not a "silver bullet"?

## Scalability issues

We need additional stuff to maintain cluster

## Complexity

When data doesn't fit table structure schema becomes too complex

**Do we have any  
alternatives?**

Do we have any  
alternatives?

The answer is  
"YES! noSQL!"

# What does CAP Theorem mean?

**2000**  
conjecture made  
by **Eric Brewer**


CA Category  
RDBMS  
GreenPlum etc

Later was proved by  
**Gilbert and Lynch**




**BASE**, an acronym for  
Basically Available Soft-  
State services with  
**Eventual Consistency**

CP Category  
BigTable  
HyperTable  
HBase  
MongoDB  
Redis  
MemCacheDB  
Scalaris etc



AP Category  
Dynamo  
CouchDB  
Cassandra  
SimpleDB  
Tokyo Cabinet  
Riak  
Voldemort etc

Stop following me, you fucking freaks!



## Storages Classification:

- **Key-Value Stores:**  
Oracle Coherence, Redis, Kyoto Cabinet etc
- **BigTable-style Databases:**  
Apache HBase, Apache Cassandra etc
- **Document Databases:**  
MongoDB, CouchDB etc
- **Full Text Search Engines:**  
Apache Lucene, Apache Solr etc
- **Graph Databases:**  
neo4j, FlockDB etc
- **RDBMS**

# Any use cases for noSQL?

- Bigness
- Massive write performance
- Fast key-value access
- Flexible schema and flexible datatypes
- Schema migration
- Write availability
- Easier maintainability, administration and operations
- No single point of failure
- Generally available parallel computing
- Programmer ease of use
- Use the right data model for the right problem
- Avoid hitting the wall
- Distributed systems support
- Tunable CAP tradeoffs



# What was our use case?

**In our case:**

- **all our data was stored in HDFS**
- **expected amount of data was several Tb**
- **application was going to build reports via Hadoop MapReduce jobs**
- **reports should stored in prepared state**
- **sorting for every column must be available**
- **paging also should be in place**

# HBase

- **Written in:** Java
- **Main point:** Billions of rows X millions of columns
- **License:** Apache
- **Protocol:** HTTP/REST (also Thrift)
- Modeled after Google's BigTable
- Uses Hadoop's HDFS as storage
- Map/reduce with Hadoop
- Query predicate push down via server side scan and get filters
- Optimizations for real time queries
- A high performance **Thrift** gateway
- HTTP supports XML, Protobuf, and binary
- Cascading, hive, and pig source and sink modules
- Jruby-based (JIRB) shell
- Rolling restart for configuration changes and minor upgrades
- Random access performance is like MySQL

**Best used:** When you use the Hadoop/HDFS stack. When you need random, realtime read/write access to BigTable-like data.

**For example:** For data that's similar to a search engine's data

# HBase

row	column families	
	info:	course:
<student_id>	info:name info:sex info:age	course:<course_id>=type

row	column families	
	info:	student:
<course_id>	info:title info:introduction info:teacher_id	student:<student_id>=type

# HBase

## **CREATE** new Table:

```
Configuration configuration = getHBaseConfiguration();
HBaseAdmin hBaseAdmin = new HBaseAdmin(configuration);

HTableDescriptor tableDescriptor = new HTableDescriptor
("students");

HColumnDescriptor info= new HColumnDescriptor("info");
info.setCompressionType(Compression.Algorithm.GZ);

tableDescriptor.addFamily(info);
....

hBaseAdmin.createTable(tableDescriptor);
hBaseAdmin.close();
```

# HBase

```
HTable table = new HTable(getHBaseConfiguration(), "students");
```

```
long studentId = 1;
```

```
byte[] rowKey = Bytes.toBytes(studentId);
```

## **//Insert data:**

```
Put value = new Put(rowKey);
```

```
value.add(Bytes.toBytes("info"), Bytes.toBytes("name")
```

```
        Bytes.toBytes("John"));
```

```
....
```

```
table.put(value);
```

## **//Get one row by key:**

```
Get get = new Get(rowKey);
```

```
table.get(get);
```

# HBase

```
HTable table = new HTable(getHBaseConfiguration(), "students");
```

```
long student1Id = 1;
```

```
long student2Id = 100;
```

```
byte[] startKey= Bytes.toBytes(student1Id);
```

```
byte[] endKey= Bytes.toBytes(student2Id);
```

## **//Get data range:**

```
Scan scan = new Scan();
```

```
scan.setStartRow(startKey);
```

```
scan.setStopRow(endKey);
```

```
ResultScanner scanner = table.getScanner(scan);
```

```
for (Result result = scanner.next();
```

```
    result != null;
```

```
    result = scanner.next())
```

```
{
```

```
    DomainObject resultObject = processResult(result);
```

```
}
```

# HBase

Problems in our case:

- HBase doesn't have support for indexes
- Index emulation causes data duplication
- Sorting and Paging was not efficient enough

# MongoDB

- Open Source.
- High performance.
- Schema free.
- Document oriented.

**Past:** part of a cloud web development platform.

**Now:** separate database project.

- Written in C++.
- Lots of language drivers available. (*BSON*)



# MongoDB

## BSON:

**Binary** JSON -> Light weight+Efficient.

Binary **JSON** -> Language independent.

Easy manipulation.

Additional data types.

Fast scan-ability.

# MongoDB

## DOCUMENT EXAMPLE:

```
{
  "_id": ObjectId("4eaf9884bb6c6747b4ba5dfb"),
  "etailer_id": 7,
  "cacheServerHost": "...",
  "localUrls": {
    "0": "...1.html.gz"
  },
  "pageInfo": {
    "STORE_PRICE": "356.66",
    "MFGR_NAME": "CANON",
    "AVL_DESC": "In Stock.",
    "PROD_ID": "SX30IS",
    "SKU": "B0041RSPR8",
    "PROD_NAME": "Canon SX30IS 14.1MP Digital Camera",
    "SHIPPING_COST": "0.00",
    "MFGR_NAME_ORIGINAL": "Canon"
  }
}
```

# MongoDB

## DEMO Time

# Links

## **NoSQL storages (currently 122+):**

<http://nosql-databases.org>

<http://kkovacs.eu/cassandra-vs-mongodb-vs-couchdb-vs-redis>

## **NoSQL modeling techniques:**

<http://highlyscalable.wordpress.com/2012/03/01/nosql-data-modeling-techniques/>

## **CAP theorem, data integrity:**

<http://www.julianbrowne.com/article/viewer/brewers-cap-theorem>

<http://www.scribd.com/doc/50353861/NoSQL-Availability-amp-Integrity-2>